APACHE PIG

Amazon Daily Review Count

data1 = load '/user/hadoop/AmazonReviews.tsv' using PigStorage('\t') AS (marketplace, customer_id, review_id, product_id, product_parent, product_title, product_category, star_rating, helpful_votes, total_votes, vine, verified_purchase, review_headline, review_body, review_date);

data = STREAM data1 THROUGH `tail -n +2` AS (marketplace, customer_id, review_id, product_id, product_parent, product_title, product_category, star_rating, helpful_votes, total_votes, vine, verified_purchase, review_headline, review_body, review_date);

daily = GROUP data by review_date;

daily_reviews = FOREACH daily GENERATE group as review_date, COUNT(data.review_id) as count;
order_by_data = ORDER daily_reviews BY count DESC;

store order_by_data INTO '/pig1';

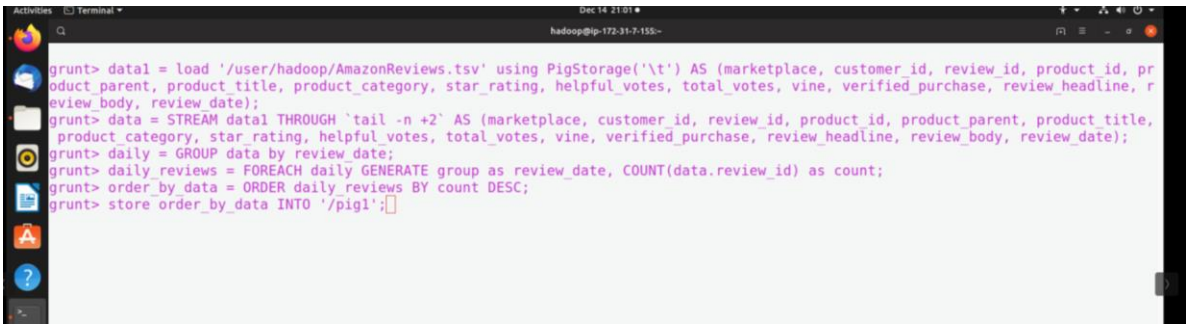

OUTPUT:


Activities
Terminal
Dec 14 21:30
hadoop@ip-172-31-7-155:~
[hadoop@ip-172-31-7-155 ~]$ hadoop fs -cat /pig1/part-v004-o000-r-00000 | head
2015-06-03 10599
2015-03-10 8641
2015-07-22 8596
2015-05-11 8568
2015-06-23 8364
2015-08-04 8315
2015-04-20 8179
2015-08-31 8114
2015-03-24 8109
2015-03-13 8100
[hadoop@ip-172-31-7-155 ~]$

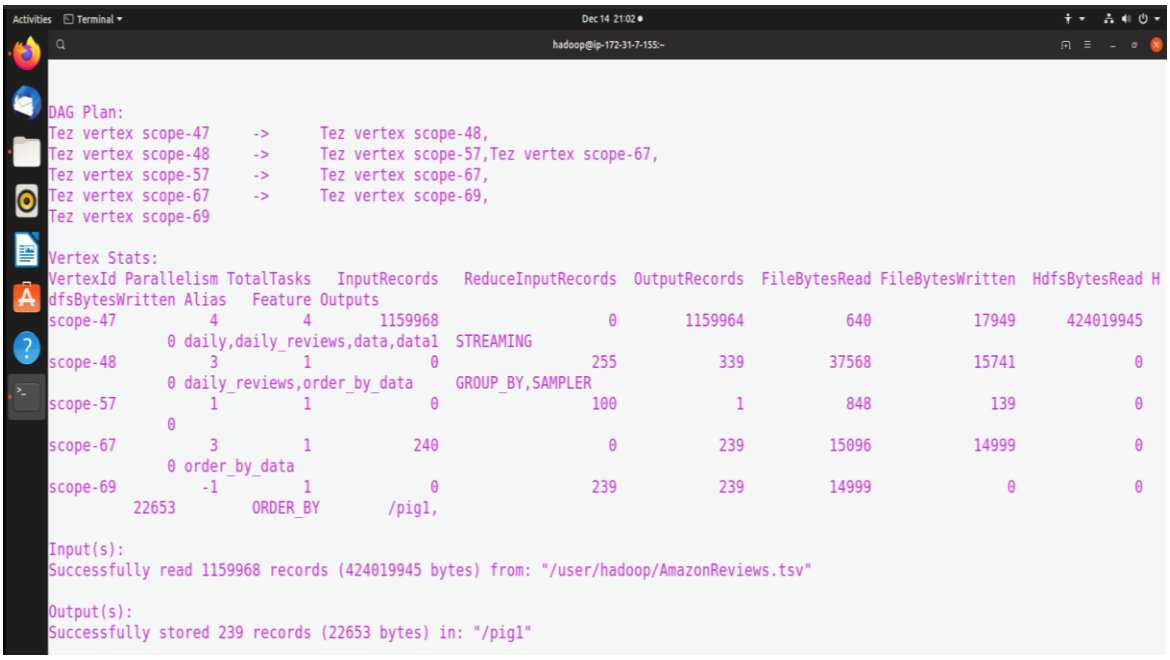
Activities
Terminal
Dec 14 21:02
hadoop@ip-172-31-7-155:~
DAG Plan:
Tez vertex scope-47 -> Tez vertex scope-48,
Tez vertex scope-48 -> Tez vertex scope-57,Tez vertex scope-67,
Tez vertex scope-57 -> Tez vertex scope-67,
Tez vertex scope-67 -> Tez vertex scope-69,
Tez vertex scope-69
Vertex Stats:
VertexId Parallelism TotalTasks InputRecords ReduceInputRecords OutputRecords FileBytesRead FileBytesWritten HdfsBytesRead H
dfsBytesWritten Alias Feature Outputs
scope-47 4 4 1159968 0 1159964 640 17949 424019945
0 daily,daily_reviews,data,data1 STREAMING
scope-48 3 1 0 255 339 37568 15741 0
0 daily_reviews,order_by_data GROUP_BY,SAMPLER
scope-57 1 1 0 100 1 848 139 0
0
scope-67 3 1 240 0 239 15096 14999 0
0 order_by_data
scope-69 -1 1 0 239 239 14999 0 0
22653 ORDER_BY /pig1,
Input(s):
Successfully read 1159968 records (424019945 bytes) from: "/user/hadoop/AmazonReviews.tsv"
Output(s):
Successfully stored 239 records (22653 bytes) in: "/pig1"

Amazon Total review count per product

data1 = load '/user/hadoop/AmazonReviews.tsv' using PigStorage('\t') AS (marketplace, customer_id, review_id, product_id, product_parent, product_title, product_category, star_rating, helpful_votes, total_votes, vine, verified_purchase, review_headline, review_body, review_date);

data = STREAM data1 THROUGH `tail -n +2` AS (marketplace, customer_id, review_id, product_id, product_parent, product_title, product_category, star_rating, helpful_votes, total_votes, vine, verified_purchase, review_headline, review_body, review_date);

prod = GROUP data by star_rating;

prod_count = FOREACH prod GENERATE group as star_rating, COUNT(data.product_id) as count;

store prod_count INTO '/pig2';

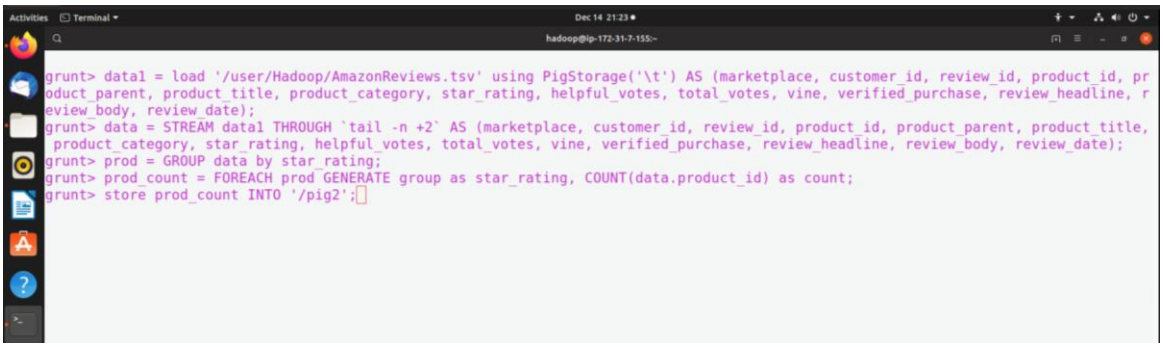

OUTPUT :

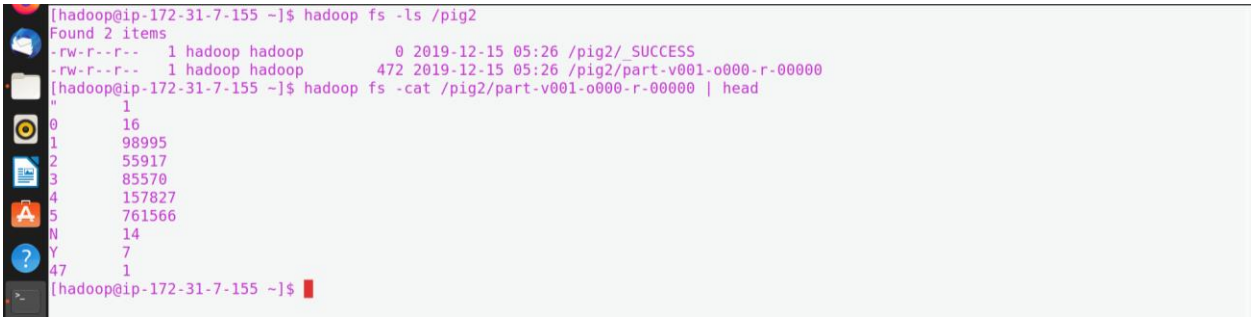

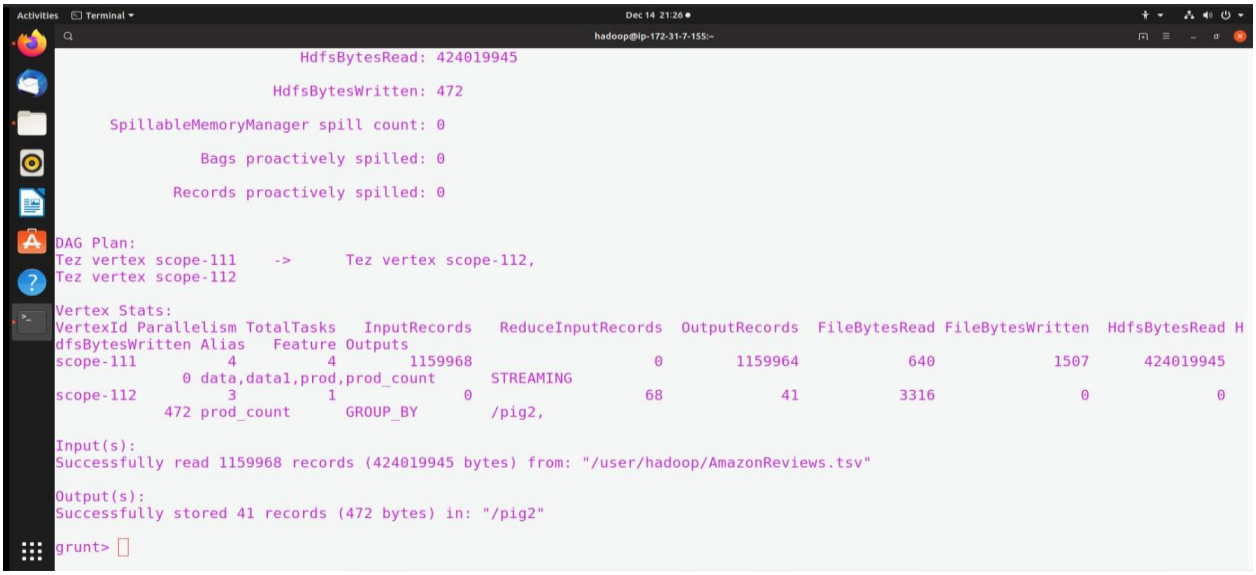
Activities
Terminal
Dec 14 21:26
hadoop@ip-172-31-7-155:~
HdfsBytesRead: 424019945
HdfsBytesWritten: 472
SpillableMemoryManager spill count: 0
Bags proactively spilled: 0
Records proactively spilled: 0
DAG Plan:
Tez vertex scope-111 -> Tez vertex scope-112,
Tez vertex scope-112
Vertex Stats:
VertexId Parallelism TotalTasks InputRecords ReduceInputRecords OutputRecords FileBytesRead FileBytesWritten HdfsBytesRead H
dfsBytesWritten Alias Feature Outputs
scope-111 4 4 1159968 0 1159964 640 1507 424019945
0 data,data1,prod,prod_count STREAMING
scope-112 3 1 0 68 41 3316 0 0
472 prod_count GROUP_BY /pig2,
Input(s):
Successfully read 1159968 records (424019945 bytes) from: "/user/hadoop/AmazonReviews.tsv"
Output(s):
Successfully stored 41 records (472 bytes) in: "/pig2"
grunt>